

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

em_data = pd.read_csv('/content/Dataset01-Employee_Attrition.csv')
```

```
em_data.head()
```

	satisfaction_level	last_evaluation	number_project	average_montly_hours	time_spend_company	Work_accident	left	promo
0	0.38	0.53	2	157	3	0	1	
1	0.80	0.86	5	262	6	0	1	
2	0.11	0.88	7	272	4	0	1	
3	0.72	0.87	5	223	5	0	1	
4	0.37	0.52	2	159	3	0	1	

Next steps: [Generate code with em_data](#) [New interactive sheet](#)

```
em_data.shape
```

```
(14999, 10)
```

```
em_data.columns
```

```
Index(['satisfaction_level', 'last_evaluation', 'number_project',
       'average_montly_hours', 'time_spend_company', 'Work_accident', 'left',
       'promotion_last_5years', 'Department', 'salary'],
      dtype='object')
```

```
em_data.dtypes
```

	0
satisfaction_level	float64
last_evaluation	float64
number_project	int64
average_montly_hours	int64
time_spend_company	int64
Work_accident	int64
left	int64
promotion_last_5years	int64
Department	object
salary	object

```
dtype: object
```

```
em_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   satisfaction_level    14999 non-null   float64
 1   last_evaluation      14999 non-null   float64
 2   number_project       14999 non-null   int64  
 3   average_montly_hours 14999 non-null   int64  
 4   time_spend_company   14999 non-null   int64  
 5   Work_accident        14999 non-null   int64  
 6   left                 14999 non-null   int64  
 7   promotion_last_5years 14999 non-null   int64  
 8   Department           14999 non-null   object 
 9   salary               14999 non-null   object 
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

```
em_data[em_data.duplicated()]
```

	satisfaction_level	last_evaluation	number_project	average_montly_hours	time_spend_company	Work_accident	left	p
396	0.46	0.57	2	139	3	0	1	
866	0.41	0.46	2	128	3	0	1	
1317	0.37	0.51	2	127	3	0	1	
1368	0.41	0.52	2	132	3	0	1	
1461	0.42	0.53	2	142	3	0	1	
...
14994	0.40	0.57	2	151	3	0	1	
14995	0.37	0.48	2	160	3	0	1	
14996	0.37	0.53	2	143	3	0	1	
14997	0.11	0.96	6	280	4	0	1	
14998	0.37	0.52	2	158	3	0	1	

3008 rows × 10 columns

✓ Dropping Duplicate Records

```
em_data1 = em_data.drop_duplicates()
```

```
em_data1.shape
```

```
(11991, 10)
```

```
em_data1.isnull().sum()
```

	0
satisfaction_level	0
last_evaluation	0
number_project	0
average_montly_hours	0
time_spend_company	0
Work_accident	0
left	0
promotion_last_5years	0
Department	0
salary	0

```
dtype: int64
```

✓ Data Visualisation and exploration

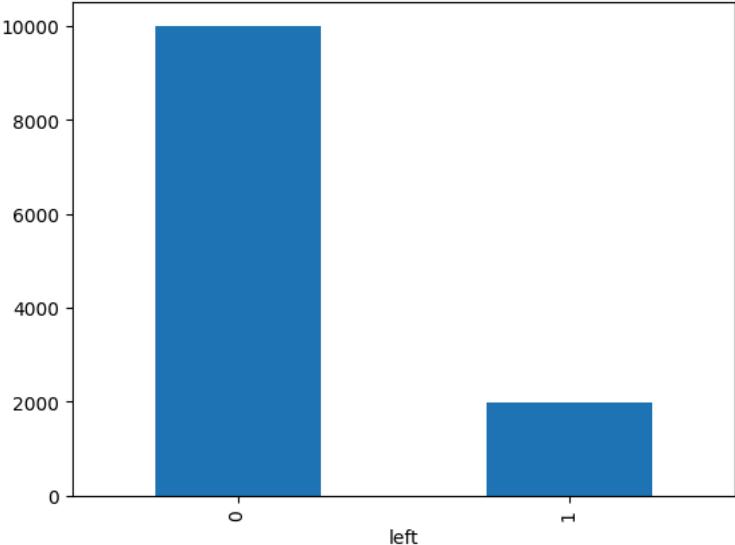
```
em_data1['left'].value_counts()
```

	count
left	
0	10000
1	1991

```
dtype: int64
```

```
em_data1['left'].value_counts().plot(kind = 'bar')
```

<Axes: xlabel='left'>



Imbalance data

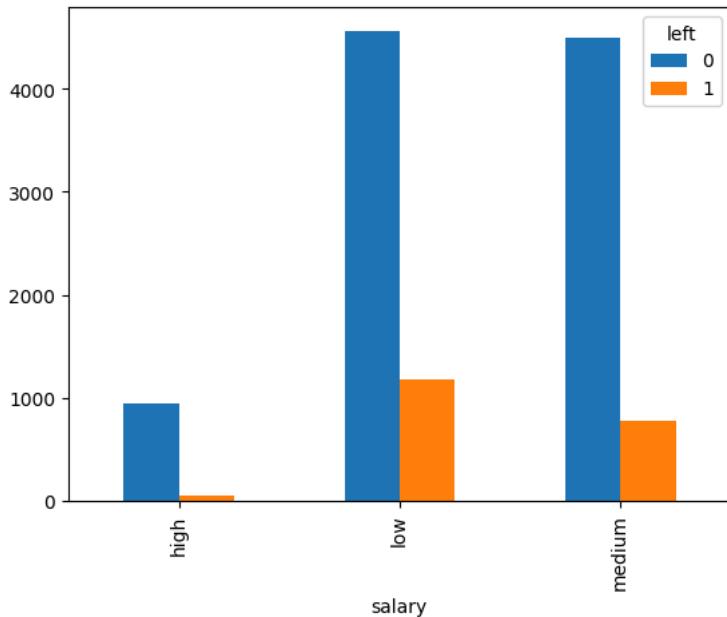
checking impact of salary on emp retention

pd.crosstab(em_data1.salary , em_data1.left)

left	0	1	
salary			grid
high	942	48	
low	4566	1174	
medium	4492	769	

pd.crosstab(em_data1.salary , em_data1.left).plot(kind = 'bar')

<Axes: xlabel='salary'>



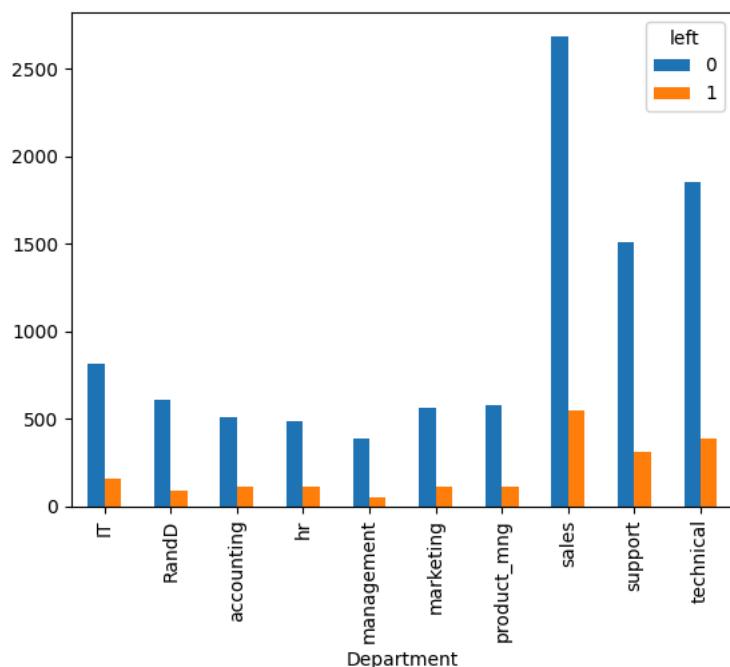
Department wise emp retention

pd.crosstab(em_data1.Department, em_data1.left)

	left	0	1
Department			
IT	818	158	
RandD	609	85	
accounting	512	109	
hr	488	113	
management	384	52	
marketing	561	112	
product_mng	576	110	
sales	2689	550	
support	1509	312	
technical	1854	390	

```
pd.crosstab(em_data1.Department, em_data1.left).plot(kind = 'bar')
```

<Axes: xlabel='Department'>



Distribution for each Numerical features

```
num_fea_list1 = [f for f in em_data1.columns if em_data1.dtypes[f] == 'float64']
```

```
num_fea_list1
```

```
['satisfaction_level', 'last_evaluation']
```

```
num_fea_list2 = [f for f in em_data1.columns if em_data1.dtypes[f] == 'int64']
```

```
num_fea_list2
```

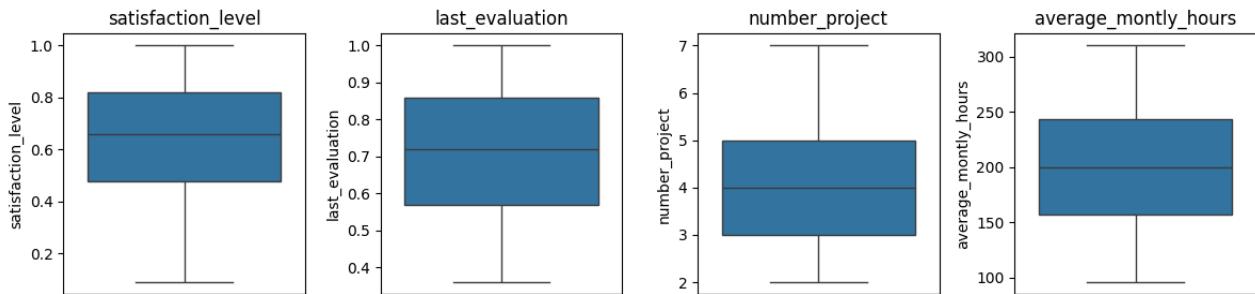
```
['number_project',
'average_montly_hours',
'time_spend_company',
'Work_accident',
'left',
'promotion_last_5years']
```

- Checking if any outliers are present or not

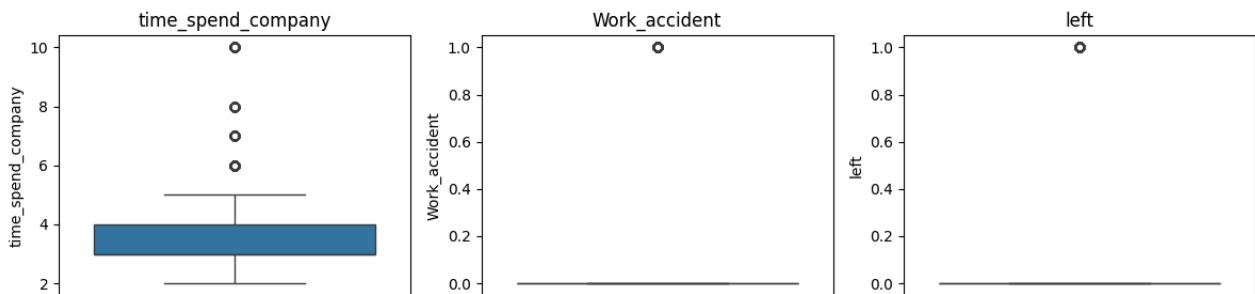
```
num_col_list = num_fea_list1 + num_fea_list2
```

```
print(num_col_list)
['satisfaction_level', 'last_evaluation', 'number_project', 'average_montly_hours', 'time_spend_company', 'Work_accident', 'left']
```

```
fig, axes = plt.subplots(ncols = 4, figsize = (12,3))
for column, axis in zip(num_col_list[:4],axes):
    sns.boxplot(data = em_data1[column], ax = axis)
    axis.set_title(column)
plt.tight_layout()
plt.show()
```



```
fig, axes = plt.subplots(ncols = 3, figsize = (12,3))
for column, axis in zip(num_col_list[4:],axes):
    sns.boxplot(data = em_data1[column], ax = axis)
    axis.set_title(column)
plt.tight_layout()
plt.show()
```



▼ Feature Engineering

Label Encoding: Converting Categorical Features to Numerical values

```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
```

```
em_data1['salary'] = label_encoder.fit_transform(em_data1['salary'])
em_data1['Department'] = label_encoder.fit_transform(em_data1['Department'])
```

```
/tmp/ipython-input-1849313355.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view
    em_data1['salary'] = label_encoder.fit_transform(em_data1['salary'])
/tmp/ipython-input-1849313355.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view
    em_data1['Department'] = label_encoder.fit_transform(em_data1['Department'])
```

```
em_data1.head()
```

	satisfaction_level	last_evaluation	number_project	average_montly_hours	time_spend_company	Work_accident	left	promo
0	0.38	0.53	2	157	3	0	1	
1	0.80	0.86	5	262	6	0	1	
2	0.11	0.88	7	272	4	0	1	
3	0.72	0.87	5	223	5	0	1	
4	0.37	0.52	2	159	3	0	1	

Next steps: [Generate code with em_data1](#) [New interactive sheet](#)

```
X = em_data1.drop('left', axis = 1)
y = em_data1['left']
```

```
from sklearn.model_selection import train_test_split
X_train,x_test , y_train, y_test = train_test_split(X, y,test_size = 0.2,random_state = 42 )
```

```
X_train.shape
```

```
(9592, 9)
```

```
from sklearn.preprocessing import StandardScaler
std_scaler = StandardScaler()
```

```
X_train_scaled = std_scaler.fit_transform(X_train)
X_test_scaled = std_scaler.fit_transform(x_test)
```

Model Development

```
from sklearn.ensemble import RandomForestClassifier
```

```
random_forest_model = RandomForestClassifier()
```

```
random_forest_model.fit(X_train_scaled, y_train)
```

```
RandomForestClassifier()
[...]
```

```
y_pred = random_forest_model.predict(X_test_scaled)
```

Model Evaluation

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[1983 15]
 [ 46 355]]
```

```
from sklearn.metrics import accuracy_score
model_accuracy = accuracy_score(y_test, y_pred)
print(f"Model accuracy : {model_accuracy}")
```

```
Model accuracy : 0.9745727386411005
```

```
from sklearn.metrics import precision_score
model_precision = precision_score(y_test, y_pred)
print(f"Model precision : {model_precision}")
```

```
Model precision : 0.9594594594594594
```

```
from sklearn.metrics import recall_score
model_recall = recall_score(y_test, y_pred)
print(f"Model recall : {model_recall}")
```

Model recall : 0.885286783042394

```
from sklearn.metrics import f1_score
model_f1_score = f1_score(y_test, y_pred)
print(f"Model f1_score : {model_f1_score}")
```

Model f1_score : 0.920881971465629

```
from sklearn.metrics import classification_report
class_report = classification_report(y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	1998
1	0.96	0.89	0.92	401
accuracy			0.97	2399
macro avg	0.97	0.94	0.95	2399
weighted avg	0.97	0.97	0.97	2399

Feature Importance

```
score_list = random_forest_model.feature_importances_
list_of_features = list(X.columns)
```

```
score_df = pd.DataFrame({'Features' : list_of_features, 'Score' : score_list})
score_df.sort_values(by = 'Score', ascending = False)
```

	Features	Score	grid
0	satisfaction_level	0.333470	grid
2	number_project	0.191268	grid
4	time_spend_company	0.164829	grid
3	average_montly_hours	0.150420	grid
1	last_evaluation	0.135869	grid
7	Department	0.011923	grid
8	salary	0.007031	grid
5	Work_accident	0.004478	grid
6	promotion_last_5years	0.000713	grid

k-fold cross validation

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(random_forest_model , X_train_scaled , y_train, cv = 5, scoring = 'accuracy')
print('Cross Validation scores = ', scores)
```

Cross Validation scores = [0.9885357 0.98332465 0.98488008 0.98488008 0.98696559]

```
avg_score = scores.mean()
```

```
print('Average Model score', avg_score)
```

Average Model score 0.9857172199849918

Hyper Parameter Tuning

```
from sklearn.model_selection import GridSearchCV
```

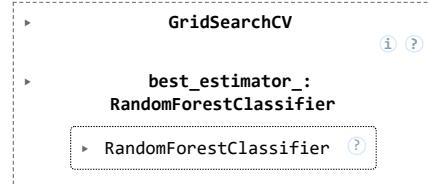
```
parameters = {
    'n_estimators' : [50,100],
    'max_features' : ['sqrt' , 'log2', None]
```

```
}
```

Searching best parameters

```
grid_search = GridSearchCV(random_forest_model, param_grid= parameters)
```

```
grid_search.fit(X_train, y_train)
```

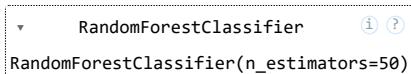


```
grid_search.best_params_
```

```
{'max_features': 'sqrt', 'n_estimators': 50}
```

```
random_forest_new_model = RandomForestClassifier(max_features= 'sqrt', n_estimators = 50)
```

```
random_forest_new_model.fit(X_train_scaled,y_train)
```



```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(random_forest_model , X_train_scaled , y_train, cv = 5, scoring = 'accuracy')
print('Cross Validation scores = ', scores)
```

```
Cross Validation scores = [0.9885357  0.98384575 0.98540146 0.98592284 0.98748697]
```

```
avg_score = scores.mean()
print('Average Model score', avg_score)
```

```
Average Model score 0.9862385420804305
```

```
Start coding or generate with AI.
```